

veth 实战

版权说明

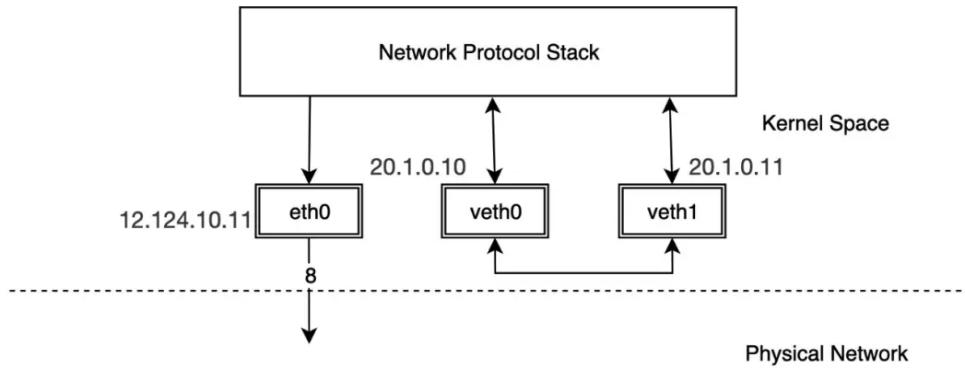
本“比特就业课”项目（以下简称“本项目”）的所有内容，包括但不限于文字、图片、音频、视频、软件、程序、数据库、设计、布局、界面等，均由本项目的开发者或授权方拥有版权。我们鼓励个人学习者使用本项目进行学习和研究。在遵守相关法律法规的前提下，个人学习者可以下载、浏览、学习本项目的内容，并为了个人学习、研究或教学目的而使用其中的材料。但请注意，未经我们明确授权，个人学习者不得将本项目的内容用于任何商业目的，包括但不限于销售、转让、许可或以其他方式从中获利。此外，个人学习者也不得擅自修改、复制、传播、展示、表演或制作本项目内容的衍生作品。任何未经授权的使用均属侵权行为，我们将依法追究法律责任。如果您希望以其他方式使用本项目的内容，包括但不限于引用、转载、摘录、改编等，请事先与我们取得联系，获取书面授权。感谢您对“比特就业课”项目的关注与支持，我们将持续努力，为您提供更好的学习体验。特此说明。比特就业课版权所有方。

对比特项目感兴趣，可以联系这个微信。



基础知识

veth pair 是成对出现的一种虚拟网络设备接口，一端连着网络协议栈，一端彼此相连。



由于它的这个特性，常常被用于构建虚拟网络拓扑。例如连接两个不同的网络命名空间(netns)，连接 docker 容器，连接网桥(Bridge)等。

常见命令

1. veth 操作

```
Shell
# 添加 veth
ip link add <veth name> type veth peer name <peer name>
# 删除 veth
ip link delete <veth name>
# 查看 veth
ip link show
```

2. 命名空间操作

```
Shell
#添加 ns
ip netns add <name>
#删除 ns
ip netns del <name>
#执行命令
ip netns exec <name> <cmd>
#遍历 ns
ip netns list
```

实战目的

了解 veth 虚拟网卡对的创建

实战步骤

1. 查看当前网卡信息

```
Shell
root@139-159-150-152:~# ifconfig
br-df863876204e: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 192.168.0.1 netmask 255.255.0.0 broadcast
192.168.255.255
        inet6 fe80::42:20ff:feb7:55bb prefixlen 64 scopeid
0x20<link>
            ether 02:42:20:b7:55:bb txqueuelen 0 (Ethernet)
            RX packets 0 bytes 0 (0.0 B)
            RX errors 0 dropped 0 overruns 0 frame 0
            TX packets 29 bytes 4216 (4.2 KB)
            TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast
172.17.255.255
        inet6 fe80::42:d7ff:fe87:d11b prefixlen 64 scopeid
0x20<link>
            ether 02:42:d7:87:d1:1b txqueuelen 0 (Ethernet)
            RX packets 586466 bytes 36737314 (36.7 MB)
            RX errors 0 dropped 0 overruns 0 frame 0
            TX packets 860320 bytes 1655179375 (1.6 GB)
            TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.0.183 netmask 255.255.255.0 broadcast
192.168.0.255
        inet6 fe80::f816:3eff:fe9d:f4ac prefixlen 64 scopeid
0x20<link>
            ether fa:16:3e:9d:f4:ac txqueuelen 1000 (Ethernet)
            RX packets 4032792 bytes 4360076332 (4.3 GB)
            RX errors 0 dropped 0 overruns 0 frame 0
            TX packets 2643585 bytes 470587289 (470.5 MB)
            TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
        loop txqueuelen 1000 (Local Loopback)
        RX packets 1126057 bytes 449407907 (449.4 MB)
```

```
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 1126057 bytes 449407907 (449.4 MB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

1xcbr0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
      inet 10.0.3.1 netmask 255.255.255.0 broadcast 10.0.3.255
      inet6 fe80::216:3eff:fe00:0 prefixlen 64 scopeid
0x20<link>
      ether 00:16:3e:00:00:00 txqueuelen 1000 (Ethernet)
      RX packets 595 bytes 72213 (72.2 KB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 817 bytes 73558 (73.5 KB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

2. 我们创建两个网络空间

```
Shell
root@139-159-150-152:~# ip netns add ns1
root@139-159-150-152:~# ip netns add ns2
root@139-159-150-152:~# ip netns list
ns2
ns1
```

3. 创建一个虚拟网卡对

```
Shell
# 创建一对 veth
ip link add veth11 type veth peer name veth12
```

4. 执行 ifconfig -a 可以看到网卡多了 2 个

```
Shell
root@139-159-150-152:~# ifconfig -a
br-df863876204e: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
      inet 192.168.0.1 netmask 255.255.0.0 broadcast
192.168.255.255
      inet6 fe80::42:20ff:feb7:55bb prefixlen 64 scopeid
0x20<link>
      ether 02:42:20:b7:55:bb txqueuelen 0 (Ethernet)
      RX packets 0 bytes 0 (0.0 B)
      RX errors 0 dropped 0 overruns 0 frame 0
```

```
TX packets 29 bytes 4216 (4.2 KB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
      inet 172.17.0.1 netmask 255.255.0.0 broadcast
172.17.255.255
      inet6 fe80::42:d7ff:fe87:d11b prefixlen 64 scopeid
0x20<link>
        ether 02:42:d7:87:d1:1b txqueuelen 0 (Ethernet)
        RX packets 586466 bytes 36737314 (36.7 MB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 860320 bytes 1655179375 (1.6 GB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
      inet 192.168.0.183 netmask 255.255.255.0 broadcast
192.168.0.255
      inet6 fe80::f816:3eff:fe9d:f4ac prefixlen 64 scopeid
0x20<link>
        ether fa:16:3e:9d:f4:ac txqueuelen 1000 (Ethernet)
        RX packets 4033449 bytes 4360150167 (4.3 GB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 2644071 bytes 470646936 (470.6 MB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
      inet 127.0.0.1 netmask 255.0.0.0
      inet6 ::1 prefixlen 128 scopeid 0x10<host>
        loop txqueuelen 1000 (Local Loopback)
        RX packets 1126165 bytes 449417393 (449.4 MB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 1126165 bytes 449417393 (449.4 MB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lxcbr0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
      inet 10.0.3.1 netmask 255.255.255.0 broadcast 10.0.3.255
      inet6 fe80::216:3eff:fe00:0 prefixlen 64 scopeid
0x20<link>
        ether 00:16:3e:00:00:00 txqueuelen 1000 (Ethernet)
        RX packets 595 bytes 72213 (72.2 KB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 817 bytes 73558 (73.5 KB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
veth11: flags=4098<Broadcast,Multicast> mtu 1500
      ether 86:f6:83:2e:31:61 txqueuelen 1000 (Ethernet)
      RX packets 0 bytes 0 (0.0 B)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 0 bytes 0 (0.0 B)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

veth12: flags=4098<Broadcast,Multicast> mtu 1500
      ether 5e:b9:7a:5d:f2:b8 txqueuelen 1000 (Ethernet)
      RX packets 0 bytes 0 (0.0 B)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 0 bytes 0 (0.0 B)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

5. 我们将网卡挪到不同的命名空间中

Shell

```
root@139-159-150-152:~# ip link set veth11 netns ns1
root@139-159-150-152:~# ip link set veth12 netns ns2
```

6. 激活我们的网卡

Shell

```
root@139-159-150-152:~# ip netns exec ns1 ip link set veth11 up
root@139-159-150-152:~# ip netns exec ns2 ip link set veth12 up
```

7. 给网卡分配 IP 地址

Shell

```
ip netns exec ns1 ip addr add 10.5.0.1/24 dev veth11
ip netns exec ns2 ip addr add 10.5.0.2/24 dev veth12
```

8. 然后我们在各自的命名空间中 ping 对方，可以看到我们的网卡是能通的

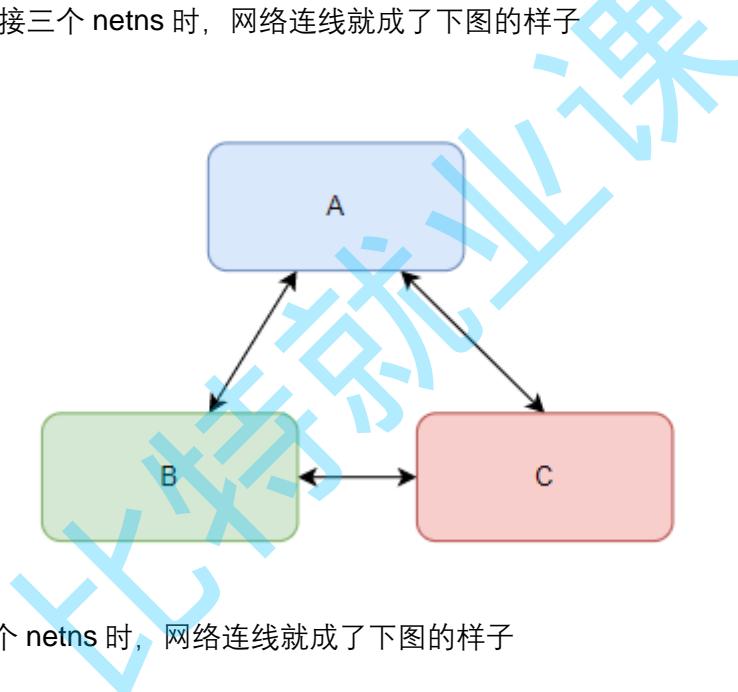
Shell

```
root@139-159-150-152:~# ip netns exec ns1 ping 10.5.0.2
PING 10.5.0.2 (10.5.0.2) 56(84) bytes of data.
64 bytes from 10.5.0.2: icmp_seq=1 ttl=64 time=0.031 ms
64 bytes from 10.5.0.2: icmp_seq=2 ttl=64 time=0.033 ms
64 bytes from 10.5.0.2: icmp_seq=3 ttl=64 time=0.033 ms
^C
--- 10.5.0.2 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2055ms
```

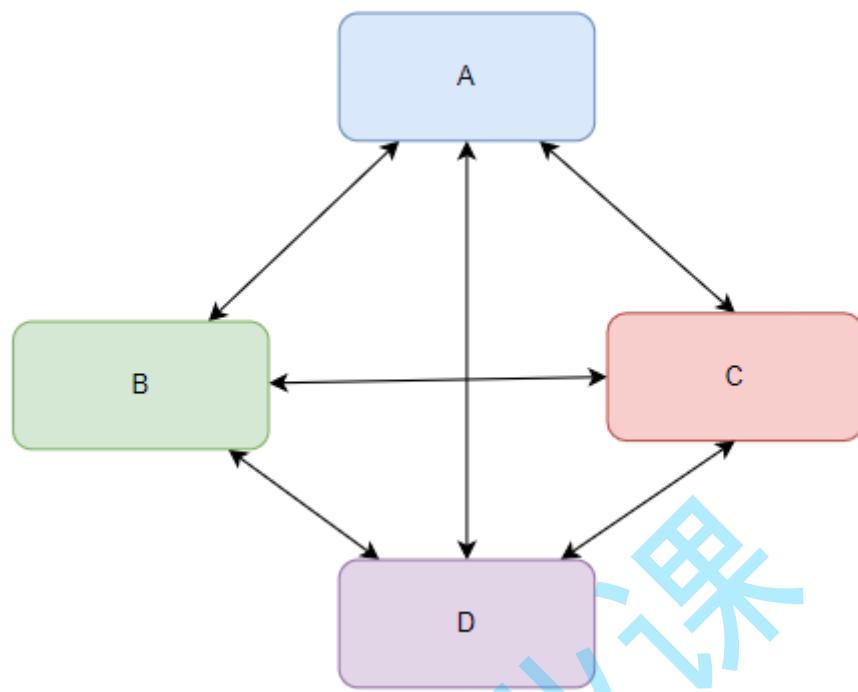
```
rtt min/avg/max/mdev = 0.031/0.032/0.033/0.000 ms
root@139-159-150-152:~# ip netns exec ns2 ping 10.5.0.1
PING 10.5.0.1 (10.5.0.1) 56(84) bytes of data.
64 bytes from 10.5.0.1: icmp_seq=1 ttl=64 time=0.023 ms
64 bytes from 10.5.0.1: icmp_seq=2 ttl=64 time=0.033 ms
64 bytes from 10.5.0.1: icmp_seq=3 ttl=64 time=0.037 ms
^X64 bytes from 10.5.0.1: icmp_seq=4 ttl=64 time=0.042 ms
^C
--- 10.5.0.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3058ms
rtt min/avg/max/mdev = 0.023/0.033/0.042/0.007 ms
```

可以看到我们使用 veth pair 将两个隔离的 netns 成功的连接到了一起。

但是这样的网络拓扑存在一个弊端，随着网络设备的增多，网络连线的复杂度将成倍增长。如果连接三个 netns 时，网络连线就成了下图的样子



而如果连接四个 netns 时，网络连线就成了下图的样子



如果有五台设备...

有没有什么技术可以解决这个问题呢？答案是有的，Linux Bridge（网桥）